



RAJESHKUMAR K

Associate DevOps Engineer · Kubernetes & Observability

+91 98946 48632 | rajeshkumar.k5445@gmail.com | India · Open to relocate
Portfolio: [rajesh-raji-98.github.io](https://github.com/rajesh-raji-98) | LinkedIn: linkedin.com/in/rajeshkumark65

PROFESSIONAL SUMMARY

DevOps engineer with 1.5+ years at Trinity Mobility, a smart-city and IoT platform company, focused on **running Kubernetes in production**. In the Technical Assistance Center (TAC) I support 20–25 live client projects: deploying releases with ArgoCD and Helm and validating the full application flow after every build, server or deployment change, resolving client issues end to end and turning each fix into a lasting improvement, and closing security gaps. I own logging and monitoring for these projects and am leading R&D on a unified observability platform with **OpenObserve** and **OpenTelemetry**.

20–25

production projects monitored

14

DLH service dashboards built

83

real-time alerts on 47 health metrics

70 / 80 / 90%

warning / critical / disaster thresholds

TECHNICAL SKILLS

Kubernetes & delivery	Kubernetes (on-prem, Azure AKS), kubectl, ArgoCD (GitOps), Helm, Calico, Docker, Nginx ingress, Linux, release and application-flow validation
Security & support	Exposed-service and Nginx hardening, certificate-expiry automation, root-cause analysis, self-healing fixes, client communication
Observability	OpenObserve, OpenTelemetry Collector (agent + gateway), Grafana, Prometheus, VictoriaMetrics, InfluxDB, ELK Stack (Filebeat, Logstash, Elasticsearch, Kibana), Alertmanager, Uptime Kuma
Tracing & telemetry	Distributed tracing, trace-log correlation (trace_id in logs), OTLP logs, Spring Boot Actuator, JMX, cAdvisor, kube-state-metrics, node / Postgres / MSSQL / MQTT / Kafka exporters
Platforms monitored	Kafka, Apache NiFi, WSO2 ESB / API Manager, Trino, Spark, Livy, Airflow, Druid, MinIO, Ranger, OpenMetadata, Hasura, MobilityDB, LDAP, Redis, Vault
Gen AI (beginner)	LiveKit and GPU monitoring, Langfuse
Scripting & query	Bash, PowerShell, Python, SQL, PromQL
Operations	Jira, severity-based alerting, SLO protection, root-cause analysis, runbooks, Linux, Git

PROFESSIONAL EXPERIENCE

Trinity Mobility Pvt Ltd — Associate DevOps Engineer

Mar 2025 – Present

Technical Assistance Center (TAC): post-go-live DevOps support and Logging & Monitoring for the company's smart-city and IoT projects

Kubernetes production support, deployments and client issues

- Support **Kubernetes-based production platforms** (on-prem and Azure AKS) for 20–25 go-live projects: pods, probes, networking (Calico), storage and ingress.
- Deploy releases with **ArgoCD and Helm**, and **validate the full application flow** (services, ingress and the real request path) after every build, server or deployment change, not just pod status.
- Own **client issues end to end**: restore service with clear updates, find every root cause, then enhance with self-healing recovery, early-warning storage alerts and automation; became a trusted point of contact for clients.
- Closed **critical security exposures** (exposed domains and services, Nginx hardening) and automated certificate-expiry tracking, ending recurring escalations.
- Documented Kubernetes fixes (for example Calico networking) as reusable runbooks for the team.

Logging & monitoring ownership

- Own logging and monitoring for **20–25 go-live projects**. Every production alert from the LnM stack raises a Jira ticket: I fix the ones in my area, and assign, follow up and close the ones that belong to other teams.

- Run severity-based alerting (**Warning 70%, Critical 80%, Disaster 90%+**) so that disk, memory, CPU and service-health problems are fixed before they turn into production events that affect the company SLO.
- Run and support the **LnM monitoring services** that every project depends on: Prometheus, VictoriaMetrics, Grafana, ELK, InfluxDB, Alertmanager (email and Telegram), Uptime Kuma, OpenTelemetry collectors, kube-state-metrics and node / Postgres / MSSQL / MQTT / Graphite exporters.
- Run the centralised **ELK logging platform** on Kubernetes, including index management, Logstash security and storage tuning for reliable log retention.

OpenObserve R&D: unified observability platform (end to end)

- Installed and built a complete **OpenObserve** stack on Kubernetes, with OpenTelemetry agent and gateway collectors (push and pull split), Kubernetes events, node and container metrics, and OTLP logs, so that every issue can be investigated in one place.
- Enabled **distributed tracing with trace-log correlation** for Java microservices (log4j MDC), WSO2 ESB and API Manager, NiFi, CEP, LDAP, IDAM and data-lake services. End-to-end trace checks passed 50/50 on the nightly environment and 14/14 on R&D.
- Built **trace-level dashboards in OpenObserve**: request rate, errors and latency per service and operation from span metrics, slowest calls, and drill-down from a panel to the trace and its correlated logs.
- Designed **45 alert pipelines and 83 real-time alerts** over 47 resource-health metrics, grouped into 9 categories and routed to Telegram, email and incident management.
- Kept the OpenObserve platform stable: diagnosed scheduler, memory and storage failures, and wrote runbooks for the team.

Data Lake House (DLH) service dashboards in Grafana (key achievement)

- Built the **OSS Engine Overview** and **13 per-service DLH dashboards** (Airflow, Trino, Spark, Livy, Kafka, Druid, MinIO, Ranger, OpenMetadata, Hasura, InfluxDB, MobilityDB and more). Each one shows UP / DOWN / UNREACHABLE health, latency, JVM and pod health on a single screen.
- Added the missing collector scrape jobs and exporter fixes that feed these dashboards, and migrated 13 OpenObserve dashboards to Grafana.

1999 vendor API: full end-to-end trace (key achievement)

- Traced the **1999 emergency-incident API** as one trace across the API gateway, WSO2 ESB, Vault, the IoT operations service and Redis, and the command-centre application, including the background fire-station recommendation flow.
- Used per-hop timings to find delays and to confirm the impact of sending the caller's reply earlier in the flow.

Gen AI observability (early work)

- Started bringing Gen AI workloads under monitoring: a LiveKit voice AI dashboard (rooms, SIP calls, call quality) and NVIDIA GPU monitoring for the GenAI servers, with Langfuse for LLM tracing.

Other contributions

- Built a port-monitoring dashboard, Airflow DAG/task dashboards and microservice health dashboards (JVM, Python, HikariCP), plus an IoT packet-journey analysis across Kafka, NiFi and CEP.
- Wrote 20+ study guides, playbooks and runbooks (OpenObserve alerting, monitoring setup, tracing, certificate fixes) that the team uses for training and incident handling.

EDUCATION

B.E. Mechanical Engineering — Kongu Engineering College, 2019
CGPA 9.01 / 10

Languages: English (fluent), Tamil (native), Hindi (basic)

CERTIFICATIONS

- Microsoft Certified: **Azure Fundamentals (AZ-900)**
- DevOps Engineer Internship — IBM with StarAgile, 2024
- NCC "C" Certificate — Cadet Under Officer, 2019